

Enterprise Computing (WS 2016/17)

Exercise 4 (3 Portfoliopunkte)

Info:

- The solution to this exercise must be handed in by Wednesday, Nov 30th 2016, 12AM.
- Any written solution must be in an accessible PDF. Any source code in a separate ZIP File. All Files are uploaded in the Information System for Instructors and Students in a single ZIP File (<https://isis.tu-berlin.de/course/view.php?id=8586>).
- Please write your name on the solution sheet.

Task 1 – Misc. (10+10+10 = 30%)

- a) Explain with your own words how distributed storage engines (like Dynamo) is handling temporary failures and give an example.
- b) Give a reasons why one should use distributed caching and explain it in a use case.
- c) Name three advantages of distributed caches and explain them in one sentence.

Task 2 – GFS (20+10+10 = 40%)

- a) Which operations are supported by GFS? For which types of queries is GFS optimized?
- b) State two main differences between Dynamo and GFS with regard to operation and data types?
- c) Explain why GFS is particularly well suited as storage system for MapReduce job input and output data (why not Dynamo or a relational database?).

Task 3 –BigTable (20+10=30%)

a) Please describe in your own words how the BigTable storage engine works by explaining the steps of a write request, and a read request, respectively (local persistence).

b) True or False: Bigtable provides additional security, geo-distribution and support for big files. If False, shortly explain.